

JOHN DOE

SYNTHETIC / FICTIONAL RESUME

Bengaluru, Karnataka, India | john.doe@example.com | +91 90000 00000 | linkedin.com/in/johndoe | github.com/johndoe

PROFESSIONAL SUMMARY

Data Engineer with 2+ years of experience designing, building, and optimizing scalable batch and streaming ETL/ELT pipelines, dimensional data warehouses, and data quality frameworks. Proficient in Python, advanced SQL, Apache Airflow, Apache Spark (PySpark), PostgreSQL, and cloud data platforms. Skilled in turning raw, unstructured data into clean, validated, analytical datasets for reporting and machine learning.

TECHNICAL SKILLS

Languages: Python, SQL, Bash

Data Processing & ETL: Apache Airflow, Apache Spark, PySpark, Pandas, NumPy

Databases & Warehouses: PostgreSQL, MySQL, Snowflake, MongoDB

Data Modeling & Concepts: Star Schema, Snowflake Schema, Data Validation, Schema Evolution, Partitioning

Tools & Cloud: Docker, Git, GitHub Actions, Linux, AWS (S3, EC2, Glue)

EXPERIENCE

Data Engineer — DataMatrix Analytics | Jul 2024 – Present

- Constructed automated ETL data pipelines in Apache Airflow and Python, processing over 2 million daily events from multiple sources.
- Designed dimensional data models and optimized PostgreSQL schemas, implementing partition pruning and indexing strategies.
- Developed Spark transformation scripts to cleanse, deduplicate, and aggregate transactional datasets into analytics-ready tables.
- Established automated data validation checks and alert notifications for pipeline failures and schema anomalies.

Associate Data Engineer — Synapse DataWorks | Jun 2023 – Jun 2024

- Maintained and scheduled batch ETL workflows extracting data from internal REST APIs into central staging databases.
- Wrote complex analytical SQL queries, stored procedures, and materialized views to support business reporting dashboards.

PROJECTS

Automated Multi-Source Data Lake Pipeline — Python, Apache Airflow, AWS S3, PySpark, PostgreSQL

- Built end-to-end data pipeline extracting public API data, landing raw files in S3, and loading transformed tables into warehouse.

E-Commerce Transaction Quality & Cleansing Engine — Python, Pandas, PostgreSQL, Docker

- Engineered automated validation system checking null thresholds, format compliance, and foreign key integrity.

EDUCATION

Bachelor of Technology in Information Technology | Anna University | 2019 – 2023 | CGPA: 8.3 / 10

CERTIFICATIONS

- IBM Data Engineering Professional Certificate
- Databricks Certified Associate Developer for Apache Spark